

Random Forest Model Training

Random forest hyperparameters tuned on full and observed datasets for each case separately

Amos Okutse

22 December, 2022

Contents

0.1	Introduction	1
0.2	Data generating process	1
0.2.1	RANDOM FOREST MODELS	2
0.3	TUNING ON THE FULL DATASET	2
0.3.0.1	Full data modelling [case 1 when $n = 500$ and $sd = 1$]	2
0.3.0.2	Full data modeling [case 2 when $n = 500$ and $sd = 45$]	5
0.3.0.3	Full data modeling [case 3 when $n = 2000$ and $sd = 1$]	8
0.3.0.4	Full data modeling [case 4 when $n = 2000$ and $sd = 45$]	11
0.4	TUNING ON OBSERVED DATA	14
0.4.0.1	[case 1 when $n = 500$ and $sd = 1$]	14
0.4.0.2	[case 2 when $n = 500$ and $sd = 45$]	17
0.4.0.3	[case 3 when $n = 2000$ and $sd = 1$]	20
0.4.0.4	[case 4 when $n = 2000$ and $sd = 45$]	23
0.5	Variable Importance	26
0.5.1	Case when $n = 500$, $SD = 1$	26
0.5.2	Case when $n = 500$, $SD = 45$	26

0.1 Introduction

Generate the data and see whether the two-step process results in less biased estimates? [We are adjusting for covariates that did not generate the outcome Y , at least not directly since the outcome was generated by Z'_i s which were not observed]

0.2 Data generating process

The data generating process is derived from Kang and Schaffer (2007) pg. 526. Assume we have a random sample, $i = 1, \dots, n$ from an infinite population. Outcome variable, y_i , R_i is the response indicator corresponding to 1 if y_i is observed and 0 otherwise. x_i is the p-dimensional vector of covariates that may be related to y_i and R_i . The population response and non-response rates are denoted by $r^{(1)} = P(R_i = 1)$ and $P(R_i = 0)$, respectively, whereas the sample response rates are denoted by $\hat{r}^{(1)} = n^{(1)}/n$ and $\hat{r}^{(0)} = n^{(0)}/n$. For each i suppose that $(z_{i1}, z_{i2}, z_{i3}, z_{i4})^T$ are independently distributed as $N(0, I)$ where I is a 4×4 identity matrix. y_i 's are generated as:

$$y_i = 210 + 50A_i + 27.4Z_{i1} + 13.7z_{i2} + 13.7z_{i3} + 13.7z_{i4} + \epsilon_i \quad (1)$$

where $\epsilon \sim N(0, 1)$ and the true propensity scores are defined as:

$$\pi_i = \text{expit}(-z_{i1} + 0.5z_{i2} - 0.25z_{i3} - 0.1z_{i4}) \quad (2)$$

The average response rate is assumed as $r^{(1)} = 0.5$ and the difference in means is $E(y/R = 1) - E(y/R = 0) = 20$. The correct π model in this scenario is a logistic regression model of R_i on the z_{ij} 's whereas the correct y model is a linear regression of y_i on the z_{ij} 's. Suppose, however, that the covariates actually observed by the data analyst are:

$$\begin{aligned} x_{i1} &= \exp(z_{i1}/2), \\ x_{i2} &= z_{i2}/(1 + \exp(Z_{i1})) + 10, \\ x_{i3} &= (z_{i1}z_{i3}/25 + 0.6)^3, \\ x_{i4} &= (z_2 + z_4 + 20)^2 \end{aligned}$$

such that $\log(\pi_i)$ and $E(y_i/x)$ are linear functions of $\log(x_{i1}), x_2, x_1^2x_2, 1/\log(x_1), x_3/\log(x_1), x_4^{1/2}$. We proceed by sampling $n = 1000$ observations as follows:

- Set $A_i = 1$ for $i = 1, \dots, 500$ and $A_i = 0$ for $i = 501, \dots, 1000$.
- Draw $z_{i1}, \dots, z_{i4} \sim N(0, I)$ where I is a 4×4 identity matrix.
- Calculate y_i as in (1).
- Calculate $\text{expit}(-z_{i1} + 0.5z_{i2} - 0.25z_{i3} - 0.1z_{i4})$
- Draw $R_i \sim \text{Ber}(\pi_i)$
- Denote $R_i = 0$ as those with missing outcome data y_i
- Calculate the x_i 's from the z_i 's

```
setwd("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\[4]_random_forest\\train_rf")
## set up the simulation
rm(list = ls())
## load the saved single data files
load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\df_one.RData")
load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\df_two.RData")
```

```
load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\df_three.RData")
load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\df_four.RData")

## load the saved list data files
##load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\dsets1.RData")
##load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\dsets2.RData")
##load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\dsets3.RData")
##load("C:\\Users\\aokutse\\OneDrive - Brown
  ↳ University\\ThesisResults\\data\\dsets4.RData")
```

0.2.1 RANDOM FOREST MODELS

0.3 TUNING ON THE FULL DATASET

0.3.0.1 Full data modelling [case 1 when $n = 500$ and $sd = 1$]

- All model parameter tuning are based on the full data.
- Parameters are explored for all possible scenarios on the full data set with 10-fold cross-validation.
- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.

```
## using data with n = 500 and sd = 1 (df_one)

train <- df_one[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =
```

```

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
    ↪ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```


## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>  <dbl> <chr>
## 1     5     4 rmse     standard  11.2   10    0.528 Preprocessor1_Model101
## 2     5     6 rmse     standard  11.4   10    0.525 Preprocessor1_Model120
## 3     4    10 rmse     standard  11.8   10    0.504 Preprocessor1_Model104
## 4     3     8 rmse     standard  11.9   10    0.472 Preprocessor1_Model125
## 5     4    12 rmse     standard  12.0   10    0.513 Preprocessor1_Model122

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
```

```
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     5     4 Preprocessor1_Model01
```

```
## variable importance??
```

```
final_md <- finalize_model(rf_model, rf_best)
options(scipen = 999)
full_one <- final_md %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = df_one[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
full_one
```




Figure 1: Variable importance for the full data case with $n = 500$ and $SD = 1$, $dpi = 300$

0.3.0.2 Full data modeling [case 2 when $n = 500$ and $sd = 45$]

- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.
- Based on this procedure, the best tuning parameters with $n = 1000$ trees were $\text{min_n} = 23$ and $\text{mtry} = 4$

```
## using data with n = 500 and sd = 45 (df_two)

train <- df_two[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)
```

```

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
      ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```


## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     4    23 rmse     standard  44.5    10   0.875 Preprocessor1_Model114
## 2     3    21 rmse     standard  44.6    10   0.911 Preprocessor1_Model111
## 3     3    27 rmse     standard  44.6    10   0.882 Preprocessor1_Model107
## 4     4    19 rmse     standard  44.6    10   0.873 Preprocessor1_Model108
## 5     3    26 rmse     standard  44.7    10   0.907 Preprocessor1_Model102

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
```

```
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     4    23 Preprocessor1_Model14
```

```
## summarize variable importance
final_md2 <- finalize_model(rf_model, rf_best)
options(scipen = 999)
full_two <- final_md2 %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = df_two[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
full_two
```





Figure 2: Variable importance for the full data case with $n = 500$ and $SD = 45$, $dpi = 300$

0.3.0.3 Full data modeling [case 3 when $n = 2000$ and $sd = 1$]

- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.
- Based on this procedure, the best tuning parameters with $n = 1000$ trees were `min_n = 4` and `mtry = 5`

```
## using data with n = 2000 and sd = 1 (df_three)

train <- df_three[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)
```

```

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
      ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```


## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     5     4 rmse     standard  8.39    10    0.254 Preprocessor1_Model101
## 2     5     6 rmse     standard  8.46    10    0.253 Preprocessor1_Model120
## 3     4    10 rmse     standard  8.78    10    0.268 Preprocessor1_Model104
## 4     4    12 rmse     standard  8.93    10    0.275 Preprocessor1_Model122
## 5     3     8 rmse     standard  8.96    10    0.298 Preprocessor1_Model125

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     5     4 Preprocessor1_Model01
```

```
## summarize variable importance
final_md3 <- finalize_model(rf_model, rf_best)
options(scipen = 999)
full_three <- final_md3 %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = df_three[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
full_three
```




Figure 3: Variable importance for the full data case with $n = 2000$ and $SD = 1$, $dpi = 300$

0.3.0.4 Full data modeling [case 4 when $n = 2000$ and $sd = 45$]

- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.
- Based on this procedure, the best tuning parameters with $n = 1000$ trees were $\text{min_n} = 35$ and $\text{mtry} = 2$

```
## using data with n = 2000 and sd = 45 (df_three)

train <- df_four[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)
```

```

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
      ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```


## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>  <dbl> <chr>
## 1     2    35 rmse     standard  46.7   10   0.637 Preprocessor1_Model123
## 2     3    38 rmse     standard  46.7   10   0.608 Preprocessor1_Model103
## 3     2    29 rmse     standard  46.7   10   0.635 Preprocessor1_Model115
## 4     3    37 rmse     standard  46.7   10   0.607 Preprocessor1_Model110
## 5     2    24 rmse     standard  46.7   10   0.632 Preprocessor1_Model118

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
```

```
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     2    35 Preprocessor1_Model23
```

```
## summarize variable importance
```

```
final_md4 <- finalize_model(rf_model, rf_best)
options(scipen = 999)
full_four <- final_md4 %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = df_four[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
full_four
```




Figure 4: Variable importance for the full data case with $n = 2000$ and $SD = 45$, dpi = 300

```
## create a grid of variable importance across all the cases under full data
↪ hyper-parameter tuning
#jpeg("full_vip.jpeg", width = 4, height = 4, units = 'in', res = 300)
#plot(x, y) # Make plot
#full_vip = grid.arrange(full_one, full_two, full_three, full_four, ncol = 2)
#full_vip
#dev.off()
```

0.4 TUNING ON OBSERVED DATA

0.4.0.1 [case 1 when $n = 500$ and $sd = 1$]

- All model parameter tuning are based on the full data.
- Parameters are explored for all possible scenarios on the full data set with 10-fold cross-validation.
- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.

```
## using data with  $n = 500$  and  $sd = 1$  (df_one)

train1 <- base::subset(df_one, R == 1)
```

```

train <- train1[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing


set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
    ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     5     4 rmse     standard  12.9   10    0.792 Preprocessor1_Model101
## 2     5     6 rmse     standard  13.1   10    0.805 Preprocessor1_Model120
## 3     3     8 rmse     standard  13.5   10    0.782 Preprocessor1_Model125
## 4     2     3 rmse     standard  13.6   10    0.724 Preprocessor1_Model113
## 5     4    10 rmse     standard  13.7   10    0.801 Preprocessor1_Model104

## can plot the best models
autoplot(rf_results)

```

```
## rf best model based on accuracy
```

```
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     5     4 Preprocessor1_Model01
```

```
## summarize variable importance
```

```
final_obs <- finalize_model(rf_model, rf_best)
options(scipen = 999)
obs_one <- final_obs %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = base::subset(df_one, R == 1)[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
obs_one
```




Figure 5: Variable importance for the observed data case with $n = 500$ and $SD = 1$, $\text{dpi} = 300$

0.4.0.2 [case 2 when $n = 500$ and $sd = 45$]

- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.
- Based on this procedure, the best tuning parameters with $n = 1000$ trees were $\text{min_n} = 23$ and $\text{mtry} = 4$

```
## using data with n = 500 and sd = 45 (df_two)

train1 <- base::subset(df_two, R == 1)
train = train[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)
```

```

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
      ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```


## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     5     4 rmse     standard  12.9   10    0.792 Preprocessor1_Model101
## 2     5     6 rmse     standard  13.1   10    0.805 Preprocessor1_Model120
## 3     3     8 rmse     standard  13.5   10    0.782 Preprocessor1_Model125
## 4     2     3 rmse     standard  13.6   10    0.724 Preprocessor1_Model113
## 5     4    10 rmse     standard  13.7   10    0.801 Preprocessor1_Model104

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     5     4 Preprocessor1_Model01
```

```
## summarize variable importance
final_obs <- finalize_model(rf_model, rf_best)
options(scipen = 999)
obs_two <- final_obs %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = base::subset(df_two, R == 1)[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
obs_two
```




Figure 6: Variable importance for the observed data case with $n = 500$ and $SD = 45$, $dpi = 300$

0.4.0.3 [case 3 when $n = 2000$ and $sd = 1$]

- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.
- Based on this procedure, the best tuning parameters with $n = 1000$ trees were $\text{min_n} = 4$ and $\text{mtry} = 5$

```
## using data with n = 2000 and sd = 1 (df_three)

train1 <- base::subset(df_three, R == 1)
train = train1[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)
```

```

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
      ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```

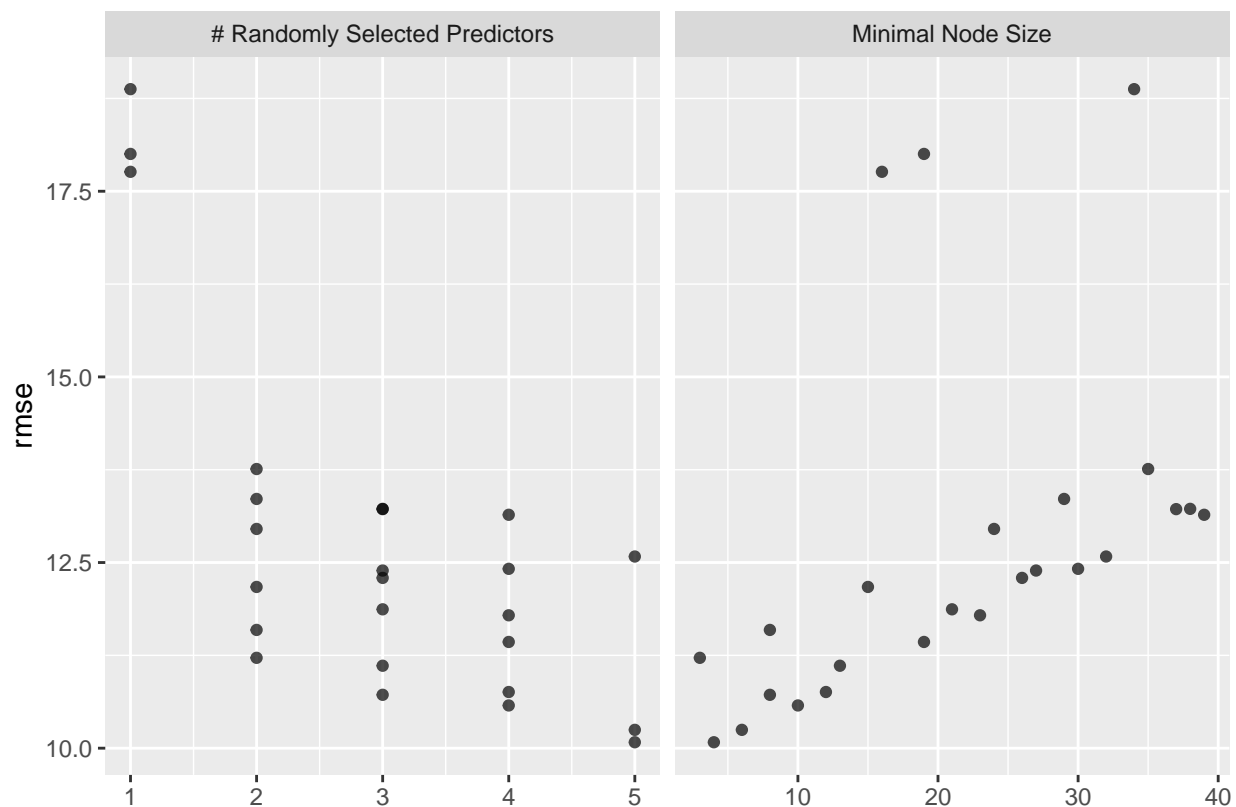
## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     5     4 rmse     standard  10.1   10    0.342 Preprocessor1_Model101
## 2     5     6 rmse     standard  10.2   10    0.330 Preprocessor1_Model120
## 3     4    10 rmse     standard  10.6   10    0.359 Preprocessor1_Model104
## 4     3     8 rmse     standard  10.7   10    0.357 Preprocessor1_Model125
## 5     4    12 rmse     standard  10.8   10    0.369 Preprocessor1_Model122

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     5     4 Preprocessor1_Model01
```

```
## summarize variable importance
final_obs <- finalize_model(rf_model, rf_best)
options(scipen = 999)
obs_three <- final_obs %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = base::subset(df_three, R == 1)[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
obs_three
```




Figure 7: Variable importance for the observed data case with $n = 2000$ and $SD = 1$, $dpi = 300$

0.4.0.4 [case 4 when $n = 2000$ and $sd = 45$]

- We use a space-filling grid design and search parameters across a grid of 25 models using cross-validation.
- Based on this procedure, the best tuning parameters with $n = 1000$ trees were $\text{min_n} = 35$ and $\text{mtry} = 2$

```
## using data with n = 2000 and sd = 45 (df_three)

train1 <- base::subset(df_four, R == 1)
train = train1[, -7]

## create the random forest model object
rf_model <- rand_forest(trees = 1000, mtry = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("regression")

## create the workflow
rf_workflow <- workflow() %>%
  add_model(rf_model) %>%
  add_formula(y ~ A + x1 + x2 + x3 + x4)
```



```

## create the procedure for validating the model
val_set <- validation_split(train, prop = 0.80)

## set up the set of metrics to gather from the models [there is no mse; can't use
  ↳ accuracy too which is for class]
metrics <- metric_set(rmse, rsq) ## rsq=coefficient of determination =

## create the re-sampling folds for hyper-parameter tuning
set.seed(345)
folds <- vfold_cv(train, v = 10)

## fit re-samples and estimate the hyper parameters
doParallel::registerDoParallel() ## leverage parallel processing

set.seed(456)
rf_results <- rf_workflow %>%
  tune_grid(val_set,
    grid = 25,
    resamples = folds,
    control = control_grid(save_pred = TRUE), #saving preds allows collecting the
      ↳ metrics
    metrics = metric_set(rmse))

rf_results %>% show_best(metric = "rmse")

```

```


## # A tibble: 5 x 8
##   mtry min_n .metric .estimator mean     n std_err .config
##   <int> <int> <chr>    <chr>    <dbl> <int>   <dbl> <chr>
## 1     3    38 rmse     standard  46.1    10    0.605 Preprocessor1_Model103
## 2     3    37 rmse     standard  46.1    10    0.599 Preprocessor1_Model110
## 3     3    26 rmse     standard  46.1    10    0.628 Preprocessor1_Model102
## 4     2    35 rmse     standard  46.1    10    0.668 Preprocessor1_Model123
## 5     4    30 rmse     standard  46.1    10    0.598 Preprocessor1_Model119

```

```

## can plot the best models
autoplot(rf_results)

```



```
## rf best model based on accuracy
```

```
rf_best = rf_results %>%
  select_best(metric = "rmse")
rf_best
```

```
## # A tibble: 1 x 3
##   mtry min_n .config
##   <int> <int> <chr>
## 1     3    38 Preprocessor1_Model103
```

```
## summarize variable importance
```

```
final_obs <- finalize_model(rf_model, rf_best)
options(scipen = 999)
obs_four <- final_obs %>%
  set_engine("ranger", importance = "impurity") %>%
  fit(y ~ .,
      data = base::subset(df_four, R == 1)[, -7]) %>%
  vip::vip(geom = "col") + theme_minimal()
obs_four
```




Figure 8: Variable importance for the observed data case with $n = 2000$ and $SD = 45$, $dpi = 300$

0.5 Variable Importance

0.5.1 Case when $n = 500$, $SD = 1$

```
#library(vip)

#final_rf %>%
#  set_engine("ranger", importance = "impurity") %>%
#  fit(y ~ .,
#    data = df_one[, -7]) %>%
#  vip(geom = "point")
```

0.5.2 Case when $n = 500$, $SD = 45$

```
#final_rf %>%
#  set_engine("ranger", importance = "impurity") %>%
#  fit(y ~ .,
#    data = df_one[, -7]) %>%
#  vip(geom = "point")
```